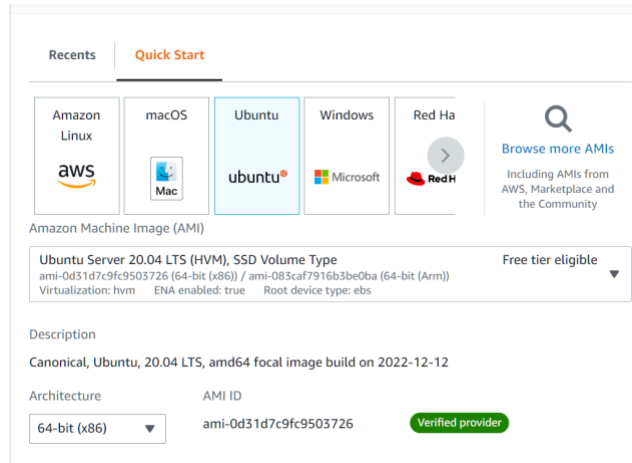


- Install EC2 instance
 - Select Ubuntu 20.04



- Create a key pair if you have not had one, download the *.pem key (make sure you remember where you put it) Note the screenshot shows dsci2024 but you can use any other name you want (e.g., I am using dsci2023).

Key pairs allow you to connect to your instance securely.

Enter the name of the key pair below. When prompted, store the private key in a secure and accessible location on your computer. **You will need it later to connect to your instance.** [Learn more](#)

Key pair name

dsci2024

The name can include upto 255 ASCII characters. It can't include leading or trailing spaces.

Key pair type

☒ RSA
RSA encrypted private and public key pair

☐ ED25519
ED25519 encrypted private and public key pair (Not supported for Windows instances)

Private key file format

☒ .pem
For use with OpenSSH

☐ .ppk
For use with PuTTY

Cancel Create key pair

- 10-20GB is sufficient (minimum is 8GB)

▼ **Configure storage**
[Info](#)

Advanced

1x
GiB
▼
Root volume (Not encrypted)

Free tier eligible customers can get up to 30 GB of EBS General Purpose (SSD) or Magnetic storage

Add new volume

The selected AMI contains more instance store volumes than the instance allows. Only the first 0 instance store volumes from the AMI will be accessible from the instance

0 x File systems
[Edit](#)

- Click Launch instance!
- Select your instance, and go to Connect. Find tab for SSH client:
- Copy the Example command
- Open a terminal with access to the ssh client program
 - If you have Windows OS, install Cygwin or use Powershell (see note at the end)
 - If you have Mac, just open a terminal windows which already has access to ssh
- “cd” to the place where you have downloaded the *.pem key.
 - Execute: `chmod 400 <your pem key>`
 - (see the AWS screenshot)
- Paste the command you copied. The command looks like this:
 - `ssh -i <your pem key> ubuntu@ec2-xxxxx.us-west-2.compute.amazonaws.com`
 - [again, see the screenshot for example](#)
 - say yes to the question.
 - You should now be connected to EC2.

Connect to your instance i-0aca171d21370d544 (Sp2023) using any of these options

EC2 Instance Connect



Session Manager

SSH client

EC2 serial console


Instance ID

 i-0aca171d21370d544 (Sp2023)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is dsci2023.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.
 `chmod 400 dsci2023.pem`
4. Connect to your instance using its Public DNS:
 `ec2-52-41-169-220.us-west-2.compute.amazonaws.com`

Example:

  Command copied `tu@ec2-52-41-169-220.us-west-2.compute.amazonaws.com`

 **Note:** In most cases, the guessed user name is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI user name.

- Note: when you restart the instance, its ip address changes. You need to recopy the ssh connection string from EC2 web site.
- Text editor on EC2:
 - `nano`
 - `vi`
- Before installing the following software, please first update package database by executing:
 - `sudo apt update`
- Install MySQL:
 - `sudo apt install mysql-server`
 - `sudo mysql`
 - In MySQL prompt (mysql>):
 - (note: do not copy and paste the following command, since it might add a new line character after '-'.)
 - `alter user 'root'@'localhost' identified with mysql_native_password by 'Dsci-551';`
 - `exit`
 - `mysql -u root -p`
 - on password prompt, type: Dsci-551 and hit enter
 - (note) MySQL server consumes a lot of main memory
 - Stop the server first, please you run other program, e.g., hdfs, spark, ...
 - Stop the server by executing:
 - `sudo service mysql stop`
 - You may start the server by executing:
 - `sudo service mysql start`
- Install Java SDK:
 - `sudo apt install default-jdk`

- (there might be a configuration menu popping up, just hit the tab key to select OK, and hit enter).
- (add this line to the end of your ~/.bashrc file on EC2)
 - export JAVA_HOME=/usr/lib/jvm/default-java
- log out and log in to EC2 again
- **note whenever you modify ~/.bashrc file, you can either execute “source ~/.bashrc” or exit EC2 and log in again so that the updated can be re-executed.**
- Install Spark:
 - (note: please install Spark, you need to install Java SDK first, see previous step)
 - wget <https://dlcdn.apache.org/spark/spark-3.5.0/spark-3.5.0-bin-hadoop3.tgz>
 - tar xvf spark-3.5.0-bin-hadoop3.tgz
 - (add this line to ~/.bashrc) <https://dlcdn.apache.org/spark/spark-3.5.1/spark-3.5.1-bin-hadoop3.tgz>
 - export PATH=\$PATH:~/spark-3.5.0-bin-hadoop3/bin
 - pyspark [spark-3.5.1-bin-hadoop3/bin](https://dlcdn.apache.org/spark/spark-3.5.1/spark-3.5.1-bin-hadoop3.tgz)
- Install Hadoop:
 - wget <https://dlcdn.apache.org/hadoop/common/hadoop-3.3.6/hadoop-3.3.6.tar.gz>
 - [tar xvf hadoop-3.3.6.tar.gz](https://dlcdn.apache.org/hadoop/common/hadoop-3.3.6/hadoop-3.3.6.tar.gz)
 - **(Skip this step if you are taking 351):** Follow the instructions in: <https://hadoop.apache.org/docs/stable/hadoop-project-dist/hadoop-common/SingleCluster.html> on “pseudo distributed operation”. In particular,
 - Follow the configuration steps
 - Follow the “set up passphraseless ssh” steps
 - Edit the file: ~/.hadoop-3.3.6/etc/hadoop/hadoop-env.sh
 - add the following line (you can edit line 54):
 - export JAVA_HOME=/usr/lib/jvm/default-java
 - follow the execution steps to format namenode, start dfs, etc.
 - add this line to ~/.bashrc file:
 - export PATH=\$PATH:~/hadoop-3.3.6/bin:~/hadoop-3.3.6/sbin
 - **note (ignore this if you are 351 students):**
 - if namenode does not start, try to reformat the namenode
 - if datanode does not come up, try:
 - rm -rf /tmp/hadoop-ubuntu/dfs/data
 - (note) this will remove the directory where hdfs stores its data node content.
 - Restart the dfs
- Install MongoDB:
 - Follow the instructions in <https://www.mongodb.com/docs/manual/tutorial/install-mongodb-on-ubuntu/>
 - Make sure using the steps for Ubuntu 20.04
 - If you get the key missing error message like this:

“W: GPG error: <https://repo.mongodb.org/apt/ubuntu> focal/mongodb-org/7.0 Release: The following signatures couldn't be verified because the public key is not available: NO_PUBKEY 160D26BB1785BA38”

Try: `sudo apt-key adv --keyserver keyserver.ubuntu.com --recv-keys 160D26BB1785BA38`

Note: replace 160D26BB1785BA38 with your missing key

- To start server: `sudo service mongod start`
- Run client: `mongosh`
- Install pip:
 - `sudo apt install python3-pip`
- Windows OS: If you are using Windows, please first download & install Cygwin.
 - If you want to use Cygwin
 - Please go to Cygwin.com
 - Download and execute `setup-x86_64.exe`
 - Make sure you select openssh package when installing
 - Your Cygwin default installation directory will be "c:\cygwin64"
 - Note: your home directory will be in msy2 will be like:
 - `c:\cygwin64\home\<your user id>`
 - copy your *.pem file downloaded from AWS to this directory