

Constraints

DSCI 551

Wensheng Wu

Kinds of Constraints

- Keys
- Foreign-key, or referential-integrity
- Value-based constraints
 - Constrain values of a particular attribute
- Tuple-based constraints
 - Constrain relationship among attributes
- Assertions: any SQL boolean expression
 - Very expressive

Keys

- Specified using "primary key" or "unique"

```
create table Sells(  
    bar varchar(100) references  
        Bars(name),  
    beer varchar(100) references  
        Beers(name),  
    price real,  
    primary key(bar, beer)  
);
```

Foreign Keys

- Consider Relation Sells(bar, beer, price).
- We might expect that a beer value is a real beer --- something appearing in Beers.name.
- A constraint that requires a beer in Sells to be a beer in Beers is called a *foreign -key* constraint.

Expressing Foreign Keys

- Use the keyword REFERENCES, either:
 1. Within the declaration of an attribute, when only one attribute is involved, or
 2. As an element of the schema, as:

FOREIGN KEY (<list of attributes>)

REFERENCES <relation> (<attributes>)

- Note MySQL seems to enforce FK only when defined as an element

Example: Express FK with Attribute

```
CREATE TABLE Beers (  
    name      CHAR(20) PRIMARY KEY,  
    manf      CHAR(20) );
```

```
CREATE TABLE Sells (  
    bar       CHAR(20) ,  
    beer      CHAR(20) REFERENCES Beers(name) ,  
    price     REAL );
```

Example: Express FK as Element

```
CREATE TABLE Beers (  
    name      CHAR(20) PRIMARY KEY,  
    manf      CHAR(20) );
```

```
CREATE TABLE Sells (  
    bar       CHAR(20),  
    beer      CHAR(20),  
    price     REAL,  
    FOREIGN KEY (beer) REFERENCES  
        Beers (name) );
```


Note parentheses are necessary!



Primary Key vs. Unique

- Referenced attributes must be declared as PRIMARY KEY or UNIQUE.
 - Otherwise, MySQL does not allow creation of the table
 - Note that primary key can not be null, but unique attribute can
- Null values can be inserted into attribute of foreign key
 - Even though it refers to primary key in referenced table


Example of FKs with Unique Attributes

- create table R (a int primary key);
 - insert into R values (1);
 - Select * from R;

Or "a int unique"
- create table S(b int, foreign key (b) references R(a));
 - insert into S values (1);
 - insert into S values (null); // this works even though "a" is primary key in R
 - select * from S;

Enforcing Foreign-Key Constraints

- If there is a foreign-key constraint from attributes of relation S to the primary key (or unique attribute) of relation R , two violations are possible:
 1. An insert or update to S introduces values not found in R .
 2. A deletion or update to R causes some tuples of S to “dangle.”



Actions Taken

- Suppose $R = \text{Beers}$, $S = \text{Sells}$.
- An insert or update to Sells that introduces a nonexistent beer must be rejected.
- A deletion or update to Beers that removes a beer value found in some tuples of Sells can be handled in three ways.



Actions Taken (Cont'd)

- The three possible ways to handle beers that suddenly cease to exist are:
 1. *Default* : Reject the modification.
 2. *Cascade* : Make the same changes in Sells.
 - Deleted beer: delete Sells tuple.
 - Updated beer: change value in Sells.
 3. *Set NULL* : Change the beers in Sells to NULL.

Example: Cascade

- Suppose we delete the Bud tuple from Beers.
 - Then delete all tuples from Sells that have beer = 'Bud'.
- Suppose we update the Bud tuple by changing 'Bud' to 'Budweiser'.
 - Then change all Sells tuples with beer = 'Bud' so that beer = 'Budweiser'.

Example: Set NULL

- Suppose we delete the Bud tuple from Beers.
 - Change all tuples of Sells that have beer = 'Bud' to have beer = NULL.
- Suppose we update the Bud tuple by changing 'Bud' to 'Budweiser'.
 - Same change.

Choosing a Policy

- When we declare a foreign key, we may choose policies SET NULL or CASCADE independently for deletions and updates.
- Follow the foreign-key declaration by:
ON [UPDATE, DELETE][SET
NULL/CASCADE]
- Two such clauses may be used.
- Otherwise, the default (reject) is used.

Example

```
CREATE TABLE Sells (  
    bar    CHAR(20),  
    beer   CHAR(20),  
    price  REAL,  
    FOREIGN KEY (beer)  
        REFERENCES Beers (name)  
    ON DELETE SET NULL  
    ON UPDATE CASCADE );
```


Multi-attribute keys (unique, PK, FK)

```
create table R(a int, b int, unique(a, b));
```

```
insert into R values(1, null);
```

```
create table P(a int, b int, primary key(a, b));
```

```
insert into P values(1, 2);
```

```
insert into R values(2, null);
```

```
create table F(a int, b int, foreign key (a, b)  
references P (a, b));
```

```
insert into F values(1, null);
```


```
create table F1 (a int, foreign key (a) references P(a));
```

```
insert into F1 values(2);
```

Are
they ok?

??

Kinds of Constraints

- Keys.
- Foreign-key, or referential-integrity.
- Value-based constraints.
 - Constrain values of a particular attribute.
- Tuple-based constraints.
 - Relationship among components.
- Assertions: any SQL boolean expression.

Attribute-Based Checks

- Put a constraint on the value of a particular attribute.
- CHECK(<condition>) must be added to the declaration for the attribute.
- The condition may use the name of the attribute, but any other relation or attribute name must be in a subquery.
- Note: MySQL does not seem to support this
 - Accept definition, but does not enforce it
- Other DBMS, e.g., PostgreSQL, may support it

Example

```
CREATE TABLE Sells (  
    bar    CHAR(20) ,  
    beer   CHAR(20)    CHECK ( beer IN  
        (SELECT name FROM Beers) ) ,  
    price  REAL CHECK ( price <= 5.00 )  
);
```

eg.

```
CREATE TABLE Employees (  
    EmployeeID INT PRIMARY KEY,  
    Name VARCHAR(50),  
    Age INT CHECK (Age >= 18 AND Age <= 65),  
    Department VARCHAR(30)  
);
```

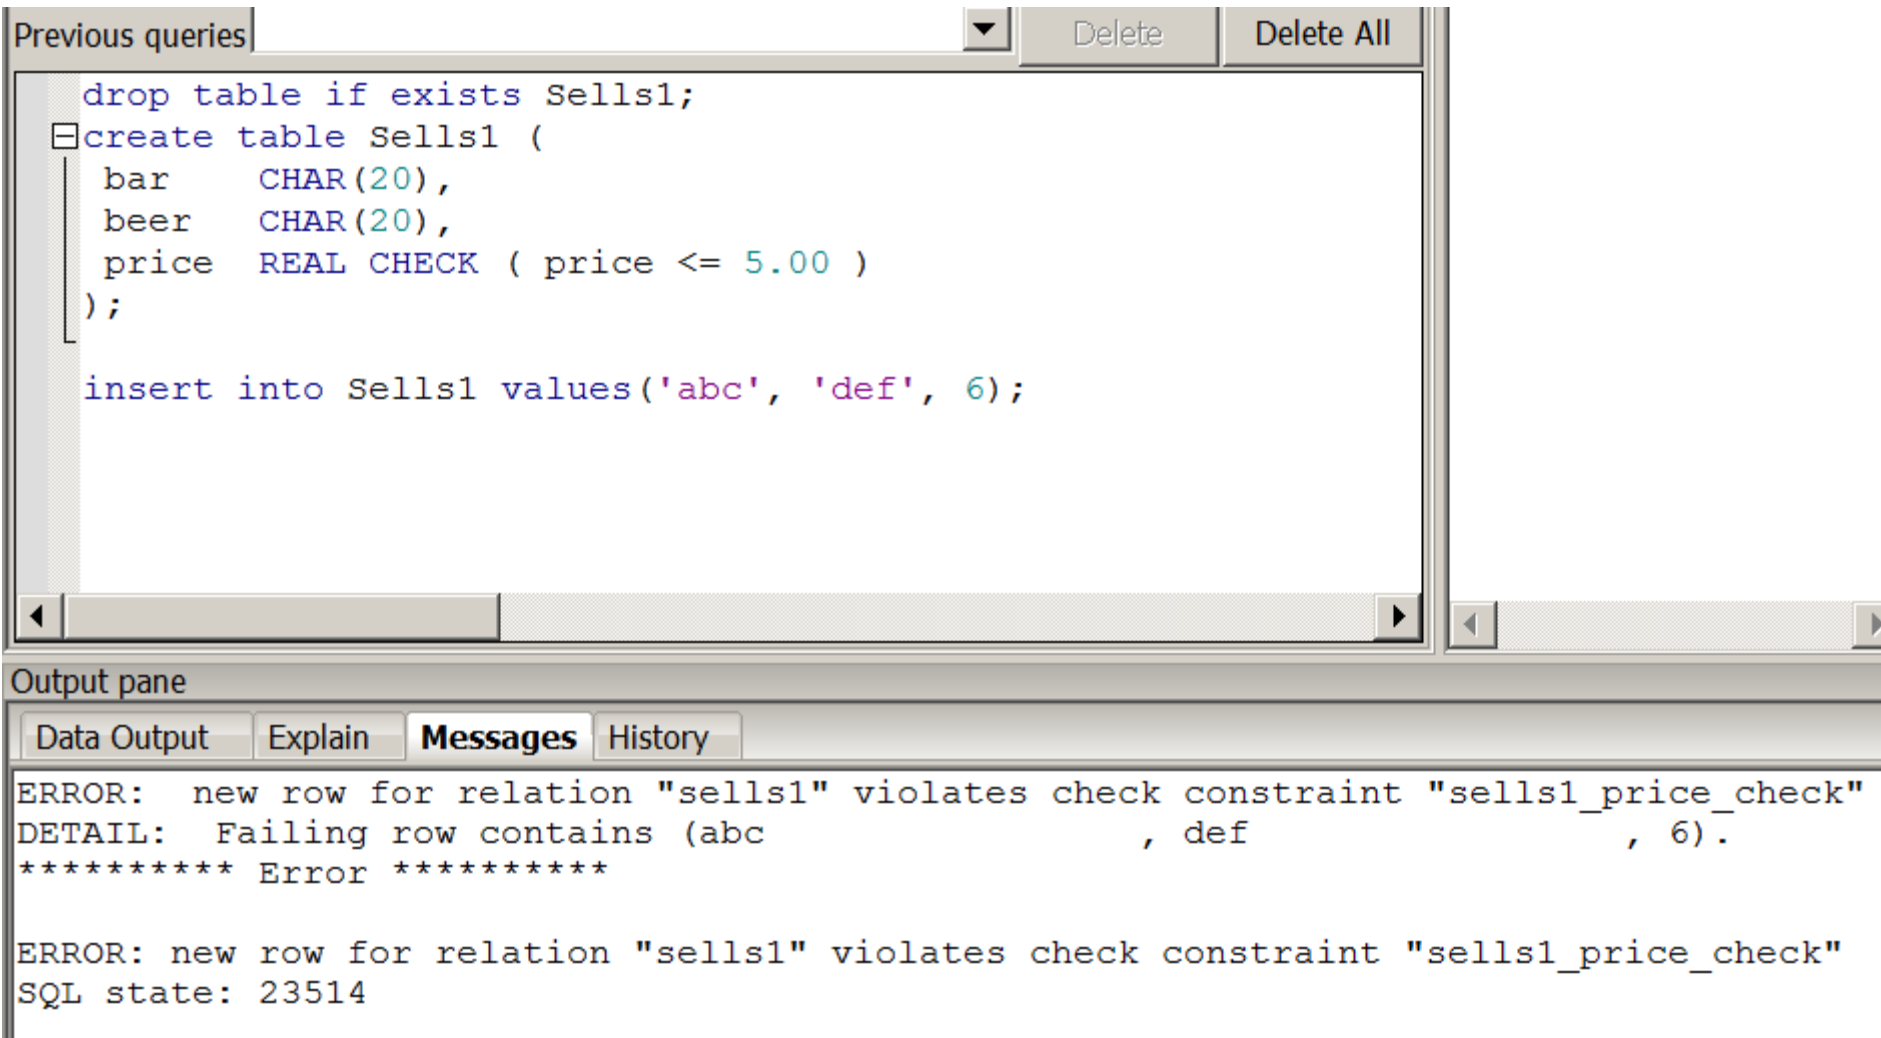
eg.

```
CREATE TABLE Customers (  
    CustomerID INT PRIMARY KEY,  
    Name VARCHAR(50),  
    Email VARCHAR(50) CHECK (Email LIKE '%@%.__%'),  
    Phone VARCHAR(20) CHECK (Phone LIKE '+__%')  
);
```

Timing of Checks

- An attribute-based check is checked only when a value for that attribute is inserted or updated.
 - Example: CHECK (price <= 5.00) checks every new price and rejects it if it is more than \$5.
 - Example: CHECK (beer IN (SELECT name FROM Beers)) not checked if a beer is deleted from Beers (unlike foreign-keys).

PostgreSQL example



The screenshot shows a PostgreSQL client window with a query editor and an output pane. The query editor contains the following SQL code:

```
drop table if exists Sells1;  
create table Sells1 (  
    bar    CHAR(20),  
    beer   CHAR(20),  
    price  REAL CHECK ( price <= 5.00 )  
);  
  
insert into Sells1 values('abc', 'def', 6);
```

The output pane shows the following error message:

```
ERROR:  new row for relation "sells1" violates check constraint "sells1_price_check"  
DETAIL:  Failing row contains (abc, def, 6).  
***** Error *****  
  
ERROR: new row for relation "sells1" violates check constraint "sells1_price_check"  
SQL state: 23514
```

Tuple-Based Checks

- CHECK (<condition>) may be added as another element of a schema definition.
- The condition may refer to any attribute of the relation, but any other attributes or relations require a subquery.
- Checked on insert or update only.

Example: Tuple-Based Check

- Only Joe's Bar can sell beer for more than \$5:

```
CREATE TABLE Sells (  
    bar          CHAR(20) ,  
    beer         CHAR(20) ,  
    price        REAL,  
    CHECK (bar = 'Joe' OR  
           price <= 5.00)  
);
```

Example (work in PostgreSQL)

- insert into sells values('Joe', 'bud', 8);
 - This insert is ok
- update sells set bar = 'joe1'
 - This update is **not** ok

Assertions

- These are database-schema elements, like relations or views.
- Defined by:

```
CREATE ASSERTION <name>  
CHECK ( <condition> );
```

- Condition may refer to any relation or attribute in the database schema.
- Very expensive to enforce
 - Neither PostgreSQL nor MySQL supports this

Example: Assertion

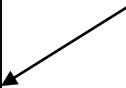
- In Sells(bar, beer, price), no bar may charge an average of more than \$5.

CREATE ASSERTION NoRipoffBars CHECK (
NOT EXISTS (

```
SELECT bar FROM Sells  
GROUP BY bar  
HAVING 5.00 < AVG(price)
```

));

Bars with an
average price
above \$5



Example: Assertion

- In Drinkers(name, addr, phone) and Bars(name, addr, license), there cannot be more bars than drinkers.

```
CREATE ASSERTION FewBar CHECK (  
    (SELECT COUNT(*) FROM Bars) <=  
    (SELECT COUNT(*) FROM Drinkers)  
);
```

Timing of Assertion Checks

- In principle, we must check every assertion after every modification to any relation of the database.
- A clever system can observe that only certain changes could cause a given assertion to be violated.
 - Example: No change to Beers can affect FewBar. Neither can an insertion to Drinkers.